

Proposal of CoNLL 2017 Shared Task: End-to-End Universal Dependency Parsing

Organizing Committee

Main contact person: Dan Zeman, zeman@ufal.mff.cuni.cz

Chair

Jan Hajič, Charles University, Prague <hajic@ufal.mff.cuni.cz>

Management Group

Joakim Nivre, Uppsala University <joakim.nivre@lingfil.uu.se>

Daniel Zeman, Charles University, Prague <zeman@ufal.mff.cuni.cz>

Filip Ginter, University of Turku <ginter@cs.utu.fi>

Slav Petrov, Google <slav@google.com>

Milan Straka, Charles University, Prague <straka@ufal.mff.cuni.cz>

Martin Popel, Charles University, Prague <popel@ufal.mff.cuni.cz>

Support Group

Christopher Manning, Stanford University

Marie-Catherine de Marneffe, Ohio State University

Yoav Goldberg, Bar Ilan University

Reut Tsarfaty, Weizmann Institute of Science

Sampo Pyysalo, University of Cambridge

Francis Tyers, UiT / The Arctic University of Norway

Jenna Kanerva, University of Turku

Lilja Øvrelid, University of Oslo

Giuseppe Celano, University of Leipzig

Miguel Ballesteros, Pompeu Fabra University

Çağrı Çöltekin, University of Tübingen

Tommi Pirinen, University of Hamburg

Paola Merlo, University of Geneva

Anssi Yli-Jyrä, University of Helsinki

Özlem Çetinoğlu, University of Stuttgart

Jon Dehdari, German Research Center for Artificial Intelligence (DFKI)

Juhani Luotolahti, University of Turku

Description of the Shared Task

Learning dependency parsers is a popular sub-task of natural language learning, useful in downstream applications as well as in linguistic research. Besides parsing proper, this task should address two related problems: working on multiple languages, including languages for which little or no data is available, and working in a real-world setting, without manually curated segmentation, tagging, etc. This task has been made possible by the Universal Dependencies initiative (UD, <http://universaldependencies.org/>), which has developed treebanks for 40+ languages with cross-linguistically consistent annotation and recoverability of the original raw texts.

Participating systems will have to find labeled syntactic dependencies between words, i.e. a syntactic *head* for each word, and a *label* classifying the type of the dependency relation. There will be multiple test sets in various languages but all data sets will adhere to the common annotation style of UD. Participants will be asked to parse raw text where no gold-standard annotation (tokenization, lemmas, morphology) is available. However, there is an open-source pipeline (UDPipe, <https://ufal.mff.cuni.cz/udpipe/>) that the participants can run instead of training their own models for any steps preceding the dependency analysis. We will even provide variants of the test data that have been preprocessed by UDPipe. We believe that this makes the task reasonably accessible for everyone.

We do not plan on running separate open and closed tracks. Instead, we want to include every system in a single track, which will be formally closed, but where the list of permitted resources is rather broad; see below for more on the data selection process.

Data

The task will only work with resources that are publicly available, royalty-free, and under a license that is free at least for non-commercial usage (e.g., CC BY-SA or CC BY-NC-SA). The right to use the data must not be limited to the shared task, i.e., the data must be available for follow-up research too.

Treebanks

The main data sets will be taken from the most recent version of the UD treebanks. Note that this means that the test data will be known in advance and we have to trust the participants not to take advantage of it. For some languages, we hope to be able to provide additional test sets that have not been previously released. These sets will not necessarily be from the same domain as the training data for the given language, but they will adhere to the UD annotation guidelines. Gold standard data from these sets will be only made available after the evaluation phase. Finally, there will be one or more **surprise language(s)**, which have not been previously released in UD and for which we will not provide training or development data, except for a small sample at the beginning of the evaluation phase. The point of having surprise languages is to encourage participants to pursue truly multilingual approaches to parsing. Gold-standard data of the surprise language(s) will be also made available after the shared task.

A conservative estimate is that we will be able to evaluate the systems on 20+ languages. We will only include UD treebanks which pass validation tests for the current UD guidelines, and for which we can obtain a test set of at least 10,000 words. There is no upper limit on the test size (the largest test set is currently ~170K). Participants will receive training+development data with gold-standard tokenization, sentence segmentation, POS tags and dependency relations; for some languages also lemmas and/or morphological features. The size of these data sets will vary according to availability. For some languages, they may be as small as the test set (or even smaller), for others it may be ten times larger than the test set, and for the surprise languages it will be close to zero. One subset of the data will be formally designated as the development set, but participants will be free to use it also for training their final system.

Raw Data

We will provide additional raw data for the languages of the shared task, useful, for example, for producing word embeddings. These data sets will be taken from CommonCrawl and automatically sorted by a language recognizer. It may not be available for all languages (note that UD contains also some classical languages such as Ancient Greek) but we are confident that we will be able to provide more than 500M words for most languages. For convenience, we will provide a variant of this data pre-processed by UDPipe, and also pre-computed word embedding vectors for those participants who want to use them but do not want to tweak their own settings of the word-to-vector software.

Parallel Data

To support multi-lingual and cross-lingual approaches and model transfers, participants are allowed to use data from the OPUS parallel corpus (<http://opus.lingfil.uu.se/>). We will not redistribute these data sets, participants are simply referred to the OPUS website.

Call for Additional Data

Instead of organizing a separate open track we will encourage the participants to report (by the end of December) additional data they want to use. If the data sets are relevant to the task and meet the public availability condition, they will be added to the list of resources available to participants.

Evaluation of Participating Systems

All systems will be required to generate valid output in the CoNLL-U format for all test sets. They will know the language of the test set, but they must respond even to unknown language codes (for which there are no training data). The systems will be able to select either raw text as input, or the file pre-processed by UDPipe. Every system must produce valid output for every test set.

The evaluation will focus on dependency relations, i.e., the index of the head node and the dependency label. POS tags, lemmas and morphological features are not considered in the main evaluation metric, although the systems are free to predict them too. On the other hand, word segmentation must be reflected in the metric because the systems do not have access to gold-standard segmentation, and identifying the words is a prerequisite for dependency evaluation.

The evaluation starts by aligning the system-produced words to the gold standard ones (see Appendix for details). Once the words are aligned, we will compute two metrics: Labeled Attachment Score (LAS) and a UD-specific metric that takes typological variation across languages into account by putting most weight on dependency relations that can be expected to be parallel across languages. Systems will be ranked by a macro-average over all test sets, and we would prefer to use the UD-specific metric for this purpose to make scores more comparable across languages. However, we still need to study the effect of adopting a new metric and will fall back on standard LAS if we are not convinced that the new metric is adequate. The final choice of the main metric should in any case be made before the first announcement of the shared task.

Labeled Attachment Score (LAS) is a standard evaluation metric in dependency parsing: the percentage of words that are assigned both the correct syntactic head and the correct dependency label. For scoring purposes, only universal dependency labels will be taken into account, which means that language-specific subtypes such as “acl:relcl” (relative clause), a subtype of the universal relation “acl” (adnominal clause), will be truncated to “acl” both in the gold standard and in the parser output in the evaluation. (Parsers can still choose to predict language-specific subtypes if it improves accuracy.) In our configuration, the standard LAS score will also have to be modified to take word segmentation mismatches into account. A dependency is therefore scored as correct only if both nodes of the relation match existing gold-standard nodes. Precision

P is the number of correct relations divided by the number of system-produced nodes; recall R is the number of correct relations divided by the number of gold-standard nodes. We define LAS as the F_1 -score = $2PR / (P+R)$.

UD-specific metric: A leading idea behind the UD representations is to focus on dependencies between content words in order to maximize parallelism across languages, because function words often correlate with morphology (or nothing at all) in other languages. Since we do not evaluate morphological analysis in this task, it makes sense to also look at results without function words. An argument for this type of metric and a first proposal can be found in <http://stp.lingfil.uu.se/~nivre/docs/udeval-cl.pdf>. Similarly to our extended LAS metric, the new metric is computed as F_1 -score of precision and recall and is sensitive to word segmentation errors. The difference to LAS is that there is a pre-defined list of content dependency relations, which are included in the evaluation or possibly weighted higher than other dependencies.

Besides the two central metrics and one overall ranking of the systems, we will evaluate the systems along various other dimensions and we may publish additional rankings for sub-tasks (e.g., performance on the surprise languages). The evaluation script is not ready at the time of writing this proposal but it will be publicly available by the end of December at the latest.

We plan to use the Tira platform (<http://www.tira.io/>) as suggested in the call for proposals. Therefore, participants will submit systems, not parsed data, allowing us to keep unreleased test data hidden until after the task has been completed.

Timeline

We agree to follow the timeline suggested in the call. Trial data (not necessarily for all languages) will be available in February 2017, train+dev in March and test data in May. No copyright-related issues have to be solved; the time between now and March will be mostly needed to refine the evaluation procedure and improve existing datasets and their compliance with the UD guidelines.

Organizing Team

The organizing team is chaired by Jan Hajič with Dan Zeman as assistant and main contact person. Together with Joakim Nivre, Filip Ginter, Slav Petrov, Milan Straka and Martin Popel, they form the management group who will do the bulk of the work. The larger support group will act as an advisory committee to the management group.

Jan Hajič chaired the organizing committee for the 2009 CoNLL shared task, and Dan Zeman co-organized two SemEval shared tasks. Joakim Nivre chaired the organizing committee for the 2007 CoNLL shared task and was a member of the committee in 2008 and 2009. The management and support groups together comprise 10 of the core members of the UD consortium and some of the leading experts on dep. parsing.

Other Relevant Information

There have been two CoNLL shared tasks in dependency parsing in 2006 and 2007. The current proposal differs from the previous tasks in several respects. We have treebanks in many languages that share the same annotation style, which makes multi-lingual approaches possible and cross-linguistic evaluation meaningful. Moreover, the number of languages is likely to exceed that of 2006 and 2007 combined. We will also propose a new evaluation metric tailored to the content-word-centric UD style. And finally, we evaluate end-to-end parsing with no gold-standard information available to the parsers.

Appendix: Data format and evaluation details

The CoNLL-U data format is described in more detail at <http://universaldependencies.org/format.html>. It is deliberately similar to the CoNLL-X format that was used in the CoNLL 2006 Shared Task and has become a de-facto standard since then. However, there are a few important extensions. Perhaps most important is the notion of *syntactic words* vs. *multi-word tokens*. It makes the tokenization step in UD harder than the relatively simple procedure called tokenization in other areas of NLP. For instance, German *zum* is a contraction of the preposition *zu* “to”, and the article *dem* “the”. In UD it is a multi-word token consisting of two syntactic words, *zu* and *dem*. These syntactic words are nodes in dependency relations. Learning this is harder than separating punctuation from words, because a contraction is not a pure concatenation of the participating words. The CoNLL-U format uses two different mechanisms here: punctuation that is conventionally written adjacent to a word is a separate single-“word” token, and an attribute in the last column tells that there was no whitespace character between the punctuation symbol and the word. On the other hand, the contraction is a multi-word token which has a separate line starting with range of following syntactic words that belong to it. Consider a German phrase *zur Stadt, zum Haus* “to the city, to the house”. The corresponding CoNLL-U section could look like this:

1-2	zur	_	_	_	_	_	_	_	_
1	zu	_	ADP	_	_	3	case	_	_
2	der	_	DET	_	_	3	det	_	_
3	Stadt	_	NOUN	_	_	0	root	_	SpaceAfter=No
4	,	_	PUNCT	_	_	3	punct	_	_
5-6	zum	_	_	_	_	_	_	_	_
5	zu	_	ADP	_	_	7	case	_	_
6	dem	_	DET	_	_	7	det	_	_
7	Haus	_	NOUN	_	_	3	conj	_	_

We will not evaluate whether the system correctly generated the range lines (1-2 *zur* and 5-6 *zum*, respectively), nor whether it generated the SpaceAfter=No attribute. But we will have to align the nodes (syntactic words) output by the system to those in the gold standard data. Thus if the system fails to recognize *zur* as a contraction and outputs

1	zur
2	Stadt
3	,

we will treat any relations going to or from the node *zur* as incorrect. The same will happen with the node “Stadt,”, should the system fail to separate punctuation from the word *Stadt*.

If the system wrongly splits the word *Haus* and outputs

7-8	Haus
7	Hau
8	das

relations involving either *Hau* or *das* will be considered incorrect.

Even if the system recognizes *zur* as contraction but outputs wrong syntactic word forms, the tokens will be considered incorrect:

1-2	zur
1	zur
2	der

Relations involving node 1 are incorrect but relations involving node 2 may be correct.

Aligning system words with the gold standard

Easy part: suppose there are no multi-word tokens (contractions). Both token sequences (gold, system) share the same underlying text (minus whitespace). Tokens can be represented as character ranges. We can find intersections of system character ranges with gold character ranges and find the alignment in one run.

Now let's assume there are multi-word tokens. They may contain anything, without any similarity to the original text; however, the data still contains the original surface form and we know to which part of the underlying text they correspond. So we only have to align the individual words between a gold and a system multi-word token. We use the LCS (longest common subsequence) algorithm for that.

Sentence boundaries will be ignored during token alignment, i.e. the entire test set will be aligned at once. The systems will have to perform sentence segmentation in order to produce valid CoNLL-U files but the sentence boundaries will be evaluated only indirectly, through dependency relations. A dependency relation that goes across a gold sentence boundary is incorrect. If on the other hand the system generates a false sentence break, it will not be penalized directly, but there must be at least one gold relation that the system did not find; not getting points for such relations will be an indirect penalization for wrong sentence segmentation.